



Trinity College Dublin

Coláiste na Tríonóide, Baile Átha Cliath

The University of Dublin

School of Computer Science and Statistics

Consumer Perception of AI-Mediated Communication in Digital Marketing Contexts: A Thematic Literature Review

Pojui Hsiao

Supervisor: Carol O'Sullivan

May 1, 2026

A dissertation submitted in partial fulfilment
of the requirements for the degree of
M.Sc. in Interactive Digital Media

Declaration

I hereby declare that this dissertation is entirely my own work and that it has not been submitted as an exercise for a degree at this or any other university.

I have read and I understand the plagiarism provisions in the General Regulations of the University Calendar for the current year, found at <http://www.tcd.ie/calendar>.

I have completed the Online Tutorial on avoiding plagiarism 'Ready Steady Write', located at <http://tcd-ie.libguides.com/plagiarism/ready-steady-write>.

I consent / do not consent to the examiner retaining a copy of the thesis beyond the examining period, should they so wish (EU GDPR May 2018).

I agree that this thesis will not be publicly available, but will be available to TCD staff and students in the University's open access institutional repository on the Trinity domain only, subject to Irish Copyright Legislation and Trinity College Library conditions of use and acknowledgement. **Please consult with your supervisor on this last item before agreeing, and delete if you do not consent**

Signed: _____

Date: _____

Abstract

This dissertation examines how artificial intelligence (AI) and generative artificial intelligence (GenAI) shape consumer perception in AI-mediated communication across digital marketing contexts. As AI becomes more visible in personalised advertising, recommendation systems, service interactions, and content creation, consumer response is shaped not only by efficiency and convenience but also by concerns about trust, authenticity, privacy, and behavioural influence. Although the literature on AI in marketing has expanded rapidly, it remains fragmented across applications and often isolates individual constructs rather than explaining how they interact.

To address this gap, the study adopts a qualitative literature review using thematic analysis. Peer-reviewed academic studies on AI, GenAI, and consumer perception in digital communication and marketing settings were selected and analysed to identify recurring themes and conceptual tensions. The review focuses on how consumers interpret AI-mediated brand communication and how these interpretations affect attitudes and behavioural intentions.

The findings show that consumer perception of AI-mediated communication is neither uniformly positive nor uniformly negative. Trust emerges as a central mediating factor, while authenticity becomes especially important when AI is involved in emotionally expressive or creative communication. Personalisation can improve perceived relevance and usefulness, but its benefits are often constrained by privacy concerns and perceived intrusiveness. Behavioural outcomes such as purchase intention, engagement, and donation intention are therefore shaped by the interaction of several perceptual mechanisms rather than by AI adoption alone.

The dissertation concludes that AI-mediated communication should be understood as a consumer-facing relational issue rather than only a technological or managerial one. It contributes to the literature by offering a consumer-centred synthesis of the main themes that shape responses to AI and GenAI in digital brand communication.

Lay Abstract

This dissertation looks at how people feel when brands and businesses use artificial intelligence to communicate with them in digital media. Today, AI is used to recommend products, personalise advertisements, answer customer questions, and even create promotional text or images. While these tools can make communication faster and more relevant, they can also make people uneasy. Some people may find AI useful and convenient, while others may feel that it is intrusive, less trustworthy, or less genuine than human communication.

The study does not collect survey or interview data. Instead, it reviews and analyses existing academic research on AI in digital marketing and brand communication. By comparing published studies, the dissertation identifies the main patterns in how consumers react to AI-generated or AI-supported communication.

The review shows that consumer reactions depend on more than just whether AI is effective. People are more likely to respond positively when AI-mediated communication feels helpful, relevant, and easy to use. However, negative reactions become more likely when AI reduces a sense of authenticity, raises privacy concerns, or makes consumers feel that communication is too automated. Overall, the dissertation argues that successful use of AI depends on how people interpret its presence in digital communication, not just on what the technology can do.

Acknowledgements

I would like to express my sincere gratitude to my supervisor for their guidance, encouragement, and feedback throughout this dissertation. Their advice helped me refine the focus of the study and develop a clearer understanding of the literature.

I am also grateful to the lecturers and staff of the School of Computer Science and Statistics for their support during my studies in Interactive Digital Media. Finally, I would like to thank my family and friends for their patience and encouragement throughout the research and writing process.

Contents

1	Introduction	1
1.1	Research Background	1
1.2	Research Problem	2
1.3	Research Aim	3
1.4	Research Objectives	3
1.5	Research Questions	3
1.6	Significance of the Study	4
1.7	Scope of the Study	4
1.8	Dissertation Structure	5
2	Literature Review	6
2.1	Introduction	6
2.2	Artificial Intelligence in Marketing	6
2.3	From Predictive AI to Generative AI	7
2.4	Consumer Perception as a Conceptual Lens	8
2.5	Key Themes in AI Marketing and Consumer Perception	10
2.5.1	Trust and AI-Enabled Customer Experience	10
2.5.2	Authenticity, Emotion, and AI-Generated Content	11
2.5.3	Personalisation, Relevance, and Perceived Usefulness	12
2.5.4	Privacy Concerns, Transparency, and Perceived Intrusiveness	12
2.5.5	Behavioural Outcomes: Purchase Intention, Engagement, and Donation Intention	13
2.6	Synthesis and Research Gap	14
2.7	Summary	14
3	Methodology	15
3.1	Research Approach	15
3.2	Research Design	15
3.3	Data Collection	16
3.4	Data Selection Criteria	16

3.5	Thematic Analysis	17
3.6	Limitations	19
3.7	Summary	19
4	Findings and Discussion	20
4.1	Introduction	20
4.2	Interpreting the Four Tensions	20
4.3	Overview of the Coding Framework	21
4.4	Data Capture and Privacy: Served versus Exploited	22
4.5	Classification and Personalisation: Understood versus Misunderstood	23
4.6	Task Delegation and Autonomy: Empowered versus Replaced	24
4.7	Social Interaction and Authenticity: Connected versus Alienated	26
4.8	How Trust, Authenticity, and Behavioural Intention Fit Across the Four Themes	27
4.9	How the Four Tensions Interact	27
4.10	Conditions That Shape More Positive or Negative Responses	28
4.11	Implications of the Findings	29
4.12	Summary	29
5	Conclusion	30
5.1	Introduction	30
5.2	Summary of Findings	30
5.3	Theoretical Implications	31
5.4	Practical Implications	32
5.5	Limitations	32
5.6	Future Research	33
5.7	Final Remarks	34
A1	Appendix	38
A1.1	Coding Framework Overview	38
A1.2	Coding Table	39
A1.3	Appendix Note	39

List of Figures

List of Tables

2.1	Operational structure of consumer perception used in this dissertation	9
A1.1	Summary of the thematic coding framework	38
A1.2	Full thematic coding table	39

1 Introduction

1.1 Research Background

Artificial intelligence (AI) is reshaping the interfaces through which brands and consumers meet in digital environments. Across social platforms, e-commerce systems, search interfaces, chatbots, and content feeds, AI increasingly helps determine how messages are produced, personalised, delivered, and interpreted. This matters not only because AI changes marketing efficiency, but because it changes the communicative form of digital brand interaction itself. Recent scholarship suggests that AI is no longer a peripheral back-end tool used only for optimisation; it increasingly participates in the design of consumer-facing communication, interaction, and media experience [1–3].

Recent industry evidence supports the scale of this transition. McKinsey reports that 88% of organisations now use AI regularly in at least one business function, with marketing and sales among the most common deployment areas, yet most firms remain in experimentation or piloting rather than full-scale transformation [4]. In parallel, Deloitte reports that GenAI is already being used in retail and consumer products for marketing and branding, audience segmentation, copy generation, sentiment analysis, and hyper-personalised shopping support, even though most organisations in these sectors remain in the early stages of adoption [5]. These figures are useful not as the main frame of the argument, but as evidence that AI-mediated communication has become a practical and visible feature of contemporary digital media environments.

The rapid diffusion of generative artificial intelligence (GenAI) has intensified this transformation. In contrast to earlier forms of predictive AI, which primarily supported classification, targeting, and forecasting, GenAI can directly produce marketing content such as text, images, product descriptions, promotional copy, and interactive brand messages. As a result, AI is no longer only involved in analysing consumer data; it is also participating in the creation of the marketing message itself [6, 7]. This development raises new questions about whether consumers interpret AI-mediated communication in the same way as human-created communication.

From an interactive digital media perspective, this shift matters because consumer responses are not determined solely by functional efficiency. Consumers also evaluate whether an AI-mediated message, interface, or interaction seems trustworthy, authentic, relevant, credible, intrusive, or helpful. These judgments shape broader outcomes such as attitudes towards the brand, willingness to engage, and purchase intention. Although AI can improve convenience and personalisation, its use may simultaneously generate unease when consumers perceive automation as impersonal, deceptive, or overly dependent on personal data [2, 8]. Understanding consumer perception is therefore central to understanding not only AI adoption in marketing, but also the changing character of digital communication and mediated interaction.

1.2 Research Problem

The use of AI in marketing presents a tension between commercial opportunity and consumer acceptance. On the one hand, AI promises improved efficiency, scalability, and relevance. Personalised recommendations, automated customer support, and AI-generated promotional content can make marketing more responsive to consumer needs and organisational goals. On the other hand, these same practices may trigger concerns about privacy, transparency, manipulation, authenticity, and the erosion of human connection in brand communication [6, 7].

This tension becomes particularly visible in studies of AI-generated and AI-disclosed marketing content. Emerging empirical work shows that consumers often respond differently when they know a message, image, or advertisement has been created by AI rather than by a human. For example, emotional marketing communication attributed to AI can reduce perceived authenticity and weaken favourable consumer responses [9]. Similarly, disclosure labels can increase novelty while also decreasing authenticity, producing mixed effects on attitudes and purchase intention [10]. Studies on AI-generated advertising and visual content also suggest that trust and purchase intention may decline when AI authorship is salient, especially in high-involvement or emotionally sensitive contexts [11–13].

At the same time, not all findings are negative. Research on AI-powered personalisation indicates that relevance, usefulness, and trust can mediate positive consumer responses, particularly when AI applications are perceived as valuable rather than merely automated [14, 15]. This suggests that consumer perception of AI marketing is neither uniformly positive nor uniformly negative. Instead, it appears to depend on a set of interconnected perceptual dimensions that operate differently across contexts, content types, and degrees of AI visibility.

Despite growing interest in this area, the literature remains fragmented. Existing studies

often examine trust, authenticity, privacy, usefulness, or purchase intention as separate issues, and many focus on specific applications such as AI-generated advertising, recommendation systems, or disclosure labels. As a result, there is still a need for a more integrated account of how AI and GenAI influence consumer perception across marketing contexts. This dissertation addresses that need by synthesising the literature through a thematic review focused explicitly on the consumer perspective.

1.3 Research Aim

The aim of this dissertation is to examine how artificial intelligence and generative artificial intelligence shape consumer perception in AI-mediated communication across digital marketing contexts.

1.4 Research Objectives

1. To review academic literature on AI and generative AI in digital marketing and brand communication.
2. To identify the main consumer perception themes that recur across this literature.
3. To analyse the mechanisms through which trust, authenticity, personalisation, privacy concerns, and perceived usefulness shape consumer responses to AI-mediated communication.
4. To explain the conditions under which AI-mediated communication is more likely to generate favourable or unfavourable consumer responses.

1.5 Research Questions

1. What are the main dimensions through which AI and generative AI shape consumer perception in digital marketing communication?
2. Through what mechanisms do trust, authenticity, personalisation, privacy concerns, and perceived usefulness influence consumer responses to AI-mediated communication?
3. Under what conditions are consumer responses to AI and generative AI in digital marketing communication more likely to become positive or negative?

1.6 Significance of the Study

This study is significant for both academic and practical reasons. Academically, it contributes to a developing body of literature at the intersection of AI, consumer behaviour, digital communication, and interactive media. While research on AI in marketing has expanded rapidly, the consumer perspective remains conceptually dispersed across different applications and theoretical lenses [3, 7]. By bringing together literature on predictive AI, generative AI, advertising, personalisation, disclosure, and mediated interaction, this dissertation offers a clearer account of the perceptual mechanisms that shape consumer responses.

Practically, the study is relevant for marketers and organisations that are increasingly incorporating AI into customer-facing activities. Businesses may assume that efficiency and innovation automatically improve consumer evaluations, yet the literature suggests that acceptance depends on how consumers interpret the meaning of AI use. If AI-generated content is seen as useful and relevant, it may strengthen engagement and purchase intention; if it is seen as inauthentic, misleading, or invasive, it may damage trust and brand relationships [9, 11, 14]. A better understanding of these dynamics can help organisations deploy AI more responsibly and strategically.

The study is also relevant to interactive digital media design and digital communication practice. As AI becomes embedded in chat interfaces, recommendation systems, disclosure cues, automated service encounters, and generative content workflows, design decisions increasingly shape how AI is experienced by users. The dissertation therefore has implications for how AI-mediated brand interactions are framed, how disclosure is designed at the user experience level, and how organisations position GenAI content within broader digital media production and communication processes.

1.7 Scope of the Study

This dissertation focuses on consumer perception in digital marketing communication rather than on the technical performance of AI systems. It examines literature on AI-enabled practices including personalised advertising, recommendation systems, AI-generated advertising content, chatbots, and social media communication. The study considers both traditional predictive AI and newer forms of generative AI, especially where AI becomes visible in the marketing message or consumer experience.

More specifically, the dissertation is concerned with consumer-AI interaction in digital media environments rather than with organisational adoption at the level of marketing strategy alone.

The dissertation is limited to secondary research and does not involve primary data collection. Its purpose is not to test causal relationships empirically, but to synthesise and interpret existing academic evidence. The emphasis is therefore on identifying recurring themes, tensions, and patterns in the literature rather than on evaluating the engineering performance of AI technologies.

1.8 Dissertation Structure

This dissertation is organised into five chapters. Chapter 1 introduces the research background, problem, aim, objectives, research questions, and scope of the study. Chapter 2 reviews the literature on AI, generative AI, and consumer perception in digital marketing communication, and develops the conceptual basis for the study. Chapter 3 explains the methodology, including the literature selection process and thematic analysis approach. Chapter 4 presents and discusses the main themes identified from the reviewed literature. Chapter 5 concludes the dissertation by summarising the findings, outlining the implications of the study, and suggesting directions for future research.

2 Literature Review

2.1 Introduction

This chapter reviews the literature on artificial intelligence (AI), generative artificial intelligence (GenAI), and consumer perception in marketing. Its main purpose is to establish the conceptual foundation for this dissertation and to identify the recurring themes that shape how consumers respond to AI-mediated marketing activity. While AI has long been used to support segmentation, targeting, forecasting, and automation, GenAI has made AI far more visible within marketing communication itself. This shift matters because consumers are no longer responding only to products, brands, or platforms. They are also responding to the role AI plays in creating, personalising, and delivering marketing messages.

Although the literature on AI in marketing has expanded rapidly, it remains somewhat fragmented. Some studies treat AI primarily as a strategic or operational capability that improves efficiency and decision-making, while others focus on consumer responses to AI-generated content, disclosure labels, recommendation systems, or personalised advertising [1, 2, 16, 17]. As a result, the field contains a wide range of useful insights, but these are often spread across different applications and theoretical traditions. This matters because the same technology may be evaluated very differently depending on whether it is used to improve convenience, automate service, or produce emotionally expressive marketing communication. This chapter therefore adopts a consumer-centred perspective and synthesises the literature around the themes most relevant to this dissertation: trust, authenticity, personalisation, privacy concerns, and behavioural outcomes.

2.2 Artificial Intelligence in Marketing

AI is now deeply embedded in contemporary marketing practice. Earlier discussions of AI in marketing often focused on its functional value, particularly its ability to analyse large datasets, predict customer behaviour, automate routine decisions, and improve targeting accuracy [1, 3]. In this earlier phase, AI largely operated in the background as a decision-support tool. Its value lay in efficiency, optimisation, and scale rather than in direct

visibility to consumers.

More recent work suggests that AI should be understood not merely as a technical tool, but as a force that is reshaping the logic of marketing itself. Davenport et al. argue that AI is changing the future of marketing by altering how firms gather customer insights, design interactions, automate communication, and make strategic decisions [16]. Kumar et al. similarly show that AI now extends across multiple marketing functions, including customer engagement, campaign design, content delivery, and decision automation [2]. In parallel, Mohd Amin et al. suggest that the adoption of AI in marketing strategies is transforming business landscapes by making predictive analytics and intelligent personalisation central to value creation [8].

Recent industry reports reinforce the practical importance of this academic discussion. McKinsey reports that AI is now regularly used by 88% of surveyed organisations in at least one business function, with marketing and sales among the most common sites of deployment and one of the main areas where firms report revenue impact [4]. Deloitte likewise shows that, in retail and consumer products, GenAI is already being applied to marketing and branding tasks such as audience segmentation, copy generation, sentiment analysis, and hyper-personalised shopping assistance [5]. At the same time, both reports stress that scaling remains difficult, governance is still evolving, and many firms continue to face concerns about inaccuracy, privacy, accountability, and workforce readiness. These industry findings do not replace academic evidence, but they help show why consumer-facing AI has become such a strategically important and contested area of marketing practice.

This broader view is important for the present dissertation because it highlights that AI does more than make marketing more efficient. It also changes how firms relate to consumers and how consumers interpret those interactions. Earlier studies are generally more optimistic when AI remains in the background as an analytical or service-support tool. By contrast, more recent work becomes more cautious when AI is visible in communication itself, because consumers begin to judge not only the outcome of marketing but also the source and meaning of the message. Once AI becomes visible within advertising, service encounters, and branded communication, consumer perception becomes a central concern rather than a secondary outcome.

2.3 From Predictive AI to Generative AI

The emergence of GenAI marks a significant shift in the literature. Hermann and Puntoni distinguish between earlier predictive forms of AI, which mainly support decisions behind the scenes, and generative systems that actively produce text, images, and other content experienced directly by consumers [6]. This distinction is particularly important in marketing

because it moves AI from an invisible analytical infrastructure to a visible participant in communication.

Chan and Choi argue that GenAI is rapidly becoming embedded across marketing practice, including advertising, service design, creative production, and innovation processes [17]. Their work is useful because it shows that GenAI should not be treated simply as an extension of traditional AI adoption. Instead, it raises new questions about content ownership, disclosure, ethics, and consumer acceptance. In a similar vein, Ubal et al. contend that consumer responses to AI depend not only on the technical capabilities of the system, but also on the psychological and contextual cues that shape whether AI is interpreted as useful, appropriate, competent, or socially acceptable [7]. A key implication here is that consumer reactions are often more positive when AI enhances speed, relevance, or accessibility, but more negative when it appears to replace human creativity or emotional judgement. For a dissertation concerned with perception, this shift from predictive AI to GenAI is especially significant because it introduces stronger tensions between efficiency, creativity, trust, and authenticity.

2.4 Consumer Perception as a Conceptual Lens

Consumer perception provides a useful lens through which to examine AI in marketing because it directs attention away from technical capability alone and towards interpretation, evaluation, and response. In this context, the central issue is not simply what AI can do, but how consumers make sense of its presence in the marketplace and how that interpretation shapes subsequent attitudes and behaviours.

For the purposes of this dissertation, consumer perception is defined as the consumer's interpretive evaluation of AI-mediated marketing encounters. This includes how consumers cognitively assess the usefulness, credibility, and appropriateness of AI in a marketing context, how they affectively respond to it, and how these evaluations shape behavioural outcomes such as engagement or purchase intention. In other words, the dissertation treats perception not as a narrow information-processing variable and not simply as a final attitude score, but as the combined process through which consumers interpret what AI is doing, what it means, and whether it feels acceptable in context.

To make this definition more operational, the dissertation distinguishes between three analytically separate levels. The first is cognitive appraisal, which includes judgements such as perceived usefulness, credibility, fairness, appropriateness, and relevance. The second is affective response, which includes trust, unease, curiosity, alienation, empathy, or moral discomfort. The third is behavioural outcome, which includes consequences such as engagement, purchase intention, ad avoidance, reliance, or resistance. In this framework,

consumer perception refers most directly to the first two levels, while behavioural outcomes are treated as downstream consequences of how consumers cognitively and affectively interpret AI-mediated encounters. This distinction is important because the literature often reports these levels together, but they should not be treated as conceptually interchangeable.

Table 2.1: Operational structure of consumer perception used in this dissertation

Level	Definition in this dissertation	Illustrative constructs from the reviewed literature
Cognitive appraisal	How consumers judge what AI is doing and whether it appears useful, credible, appropriate, fair, or intrusive	Perceived usefulness, relevance, credibility, fairness, appropriateness, opacity
Affective response	How consumers feel about the encounter once AI becomes visible or meaningful in context	Trust, unease, curiosity, alienation, empathy, moral discomfort, perceived authenticity
Behavioural outcome	What consumers become more or less likely to do after those appraisals and feelings have formed	Engagement, purchase intention, donation intention, ad avoidance, reliance, resistance

Puntoni et al. offer one of the most useful conceptual foundations for this line of inquiry. They argue that consumer experiences with AI are shaped by recurring tensions linked to data capture, classification, delegation, and social interaction [18]. These experiences may generate convenience and perceived usefulness, but they may also produce feelings of exploitation, misunderstanding, replacement, or alienation. This perspective is particularly valuable because it explains why consumer responses to AI are often mixed rather than consistently positive or negative. It also helps connect a wide range of empirical findings under a broader experiential framework.

This framework is more suitable for the present dissertation than a narrower adoption model because the dissertation is not primarily asking whether consumers intend to adopt a technology in the abstract. Nor is it focused only on a linear stimulus-response sequence. Instead, it examines how consumers interpret AI once it appears inside advertising, service, recommendation, and content-creation contexts. Puntoni et al.'s framework is especially useful because it names distinct forms of AI experience that map well onto the reviewed literature's recurring tensions. By comparison, the Technology Acceptance Model is valuable for analysing perceived usefulness and ease of use, but it is narrower and more adoption-centred than this dissertation requires. Likewise, S-O-R models are helpful for structuring stimulus, response, and organism variables, but they are less specific in

identifying the particular experiential tensions created by AI-mediated encounters. An experiential and interpretive framework is therefore more appropriate here because it allows cognitive judgements, emotional reactions, relational concerns, and behavioural consequences to be analysed together while also preserving the specificity of different AI experience domains.

Related review work reinforces this multidimensional understanding of consumer perception. Ubal et al. propose that AI-consumer interaction is influenced by psychological, contextual, and relational factors rather than by one single evaluative criterion [7]. Similarly, Hermann and Puntoni suggest that the move towards GenAI intensifies interpretive complexity because consumers are increasingly confronted with AI as a creator rather than merely a predictor [6]. Taken together, this literature suggests that trust, authenticity, personalisation, privacy, and behavioural intention should be treated as interconnected dimensions of consumer perception rather than isolated variables.

2.5 Key Themes in AI Marketing and Consumer Perception

2.5.1 Trust and AI-Enabled Customer Experience

Trust is one of the most consistently discussed themes in the AI marketing literature. It generally refers to the extent to which consumers perceive AI-enabled systems and AI-mediated messages as reliable, competent, and aligned with their interests. Trust matters because it often mediates the relationship between technological capability and favourable consumer response.

Ameen et al. make this point clearly in their study of AI-enabled customer experience. They show that trust plays a central role in shaping how consumers evaluate AI-supported service interactions, especially where convenience, personalisation, and perceived service quality are involved [19]. Their work is especially useful for this dissertation because it demonstrates that AI-driven experiences are not evaluated solely on efficiency grounds. Consumers also care about whether those experiences feel dependable, understandable, and worth the trade-offs involved.

Other studies reinforce the importance of trust in more specific marketing contexts. An and Ngo find that trust is strengthened when AI-powered personalised advertising is perceived as relevant and useful, which in turn supports purchase intention [14]. In contrast, Israfilzade shows that consumer trust can weaken when advertising is known to be AI-generated rather than human-created, especially where consumers interpret AI authorship as reducing credibility or sincerity [11]. In social media settings, Srivastava and Gurme similarly find that

trust remains a key pathway through which personalisation influences consumer behaviour, although this effect is weakened when privacy concerns become salient [15]. What emerges across these studies is a clear contextual difference: AI tends to receive more positive trust evaluations when it supports relevance and convenience behind the scenes, but more negative evaluations when it becomes highly visible in the message itself or when data use feels opaque. Overall, the literature suggests that trust is not an automatic outcome of technological sophistication; it depends on whether consumers see AI as useful, legitimate, and socially acceptable.

2.5.2 Authenticity, Emotion, and AI-Generated Content

Authenticity has become especially prominent in the GenAI literature because AI-generated content raises questions about originality, sincerity, and human presence. However, the literature does not use the term in only one way. In this dissertation, authenticity is treated as having two related but distinct dimensions. The first is source authenticity, which concerns who or what is perceived to be behind the message, such as whether the content is believed to be created by a human or by AI. The second is emotional authenticity, which concerns whether the message feels sincere, emotionally credible, or genuinely intended. This distinction is important because different studies emphasise different authenticity problems.

Arango, Singaraju, and Niininen offer strong evidence that AI-generated content can alter emotional response in meaningful ways. In their study of charitable giving advertisements, AI-generated ads reduced perceived authenticity and emotional connection, which in turn affected empathy and donation intention [20]. This is highly relevant to the present dissertation because it shows that AI-generated content may influence consumers even when the informational message remains broadly similar. What changes is the perceived genuineness of the communication and the emotional bond it creates. Their findings are most directly concerned with emotional authenticity rather than with authorship alone.

Kirk and Givi make a related argument in their discussion of the “AI-authorship effect”. They show that consumers respond less favourably to emotionally expressive marketing communication when they believe it was produced by AI rather than by a human, largely because perceived authenticity declines [9]. Here the main issue is source authenticity: the response changes because authorship is attributed to AI rather than to a human source. Shi and Jiang similarly demonstrate that AI disclosure labels can generate a mixed response: they may increase novelty, but they can also lower authenticity and weaken consumer attitudes or purchase intention [10]. These studies are broadly consistent, but they also suggest an important nuance. Negative responses become strongest in contexts where consumers expect empathy, sincerity, or human creativity, whereas novelty may partly offset

those concerns in lower-stakes or more experimental settings. Taken together, these studies suggest that authenticity should be treated as a central mechanism in consumer evaluations of GenAI marketing content rather than as a peripheral aesthetic concern, and that source authenticity and emotional authenticity should be analytically distinguished rather than merged into one undifferentiated theme.

2.5.3 Personalisation, Relevance, and Perceived Usefulness

One of the most widely cited benefits of AI in marketing is its capacity to deliver personalised experiences. AI can tailor recommendations, advertisements, and service encounters in ways that increase relevance, reduce search effort, and potentially improve customer satisfaction [1,2]. This makes personalisation a central justification for AI adoption from a managerial perspective.

However, the literature shows that personalisation is not inherently persuasive. Its positive effect depends on how consumers interpret it. Ameen et al. suggest that AI-enabled customer experience becomes more favourable when consumers perceive benefits such as convenience and improved service quality, but this evaluation is still shaped by trust and perceived sacrifice [19]. Similarly, An and Ngo find that the effect of AI-powered personalised advertising on purchase intention is strengthened when personalisation is understood as relevant and useful rather than merely technologically advanced [14]. This implies that the same mechanism can produce both positive and negative responses: personalisation is welcomed when it reduces effort and improves fit, but it becomes problematic when it feels excessive, mechanistic, or too closely tied to data surveillance. These studies suggest that personalisation works best when it feels helpful and proportionate, not when it appears intrusive or overly automated.

2.5.4 Privacy Concerns, Transparency, and Perceived Intrusiveness

Privacy concerns remain one of the most persistent tensions in the literature on AI marketing. Since many AI applications rely on the collection and interpretation of behavioural and personal data, consumers may feel uneasy about tracking, profiling, or opaque data practices. This issue becomes especially important when personalisation is perceived to cross the boundary from relevance into surveillance.

Boerman et al., although writing in the context of online behavioural advertising rather than GenAI, provide an important foundation for understanding this issue. Their review highlights how personalisation, transparency, disclosure, and perceived intrusiveness are closely linked, and it shows that consumer responses to targeted advertising are shaped by whether data

use feels visible, fair, and controllable [21]. This earlier literature remains highly relevant because AI-driven personalisation intensifies many of the same tensions.

More recent AI-focused studies confirm that privacy remains a major barrier to consumer acceptance. Davenport et al. identify privacy and bias as structural challenges in the future development of AI marketing [16]. Kumar et al. and Mohd Amin et al. likewise show that concerns around data use and ethical governance continue to shape the adoption and evaluation of AI systems in marketing [2, 8]. Empirical evidence from social media contexts also supports this view, with Srivastava and Gurme finding that privacy concerns weaken favourable behavioural responses even where personalisation and trust are otherwise present [15]. Compared with the more positive literature on relevance and convenience, this work shows that favourable responses to AI are conditional rather than stable. Privacy, then, is not just an ethical issue in the abstract; it is a perceptual barrier that can directly undermine the effectiveness of AI-enabled marketing practice.

2.5.5 Behavioural Outcomes: Purchase Intention, Engagement, and Donation Intention

Many studies ultimately connect consumer perception to behavioural outcomes such as purchase intention, engagement, recommendation, and donation intention. This is significant because it demonstrates that perceptual themes are not merely abstract attitudes. They have practical consequences for market behaviour.

Existing evidence suggests that AI can improve behavioural outcomes when consumers interpret its use as relevant, useful, and supportive of their goals. An and Ngo, for example, show that relevance, usefulness, and trust positively contribute to purchase intention in AI-powered personalised advertising [14]. However, these effects are not stable across all contexts. Arango et al. show that reduced authenticity and emotional connection can lower donation intention in charitable advertising [20]. Research on GenAI visual content similarly suggests that behavioural responses vary depending on context, involvement level, and the reasons given for AI use. Belanche et al. find that real images are often preferred in high-involvement and hedonic service contexts, while Zhang and Hur show that disclosure of motives such as privacy protection, visual appeal, or cost efficiency can alter consumer attitudes towards AI-generated advertising images [12, 13]. The pattern here is not contradictory so much as conditional: AI tends to support positive behavioural outcomes when it delivers practical value without challenging expectations of sincerity, but it becomes less effective when the product category, message type, or emotional context makes human authenticity more important. These findings indicate that behavioural outcomes are shaped by the interaction of several perceptual mechanisms rather than by AI use alone.

2.6 Synthesis and Research Gap

Taken together, the literature suggests that AI in marketing is best understood through a series of interrelated perceptual tensions. AI can improve convenience, relevance, and efficiency, yet it may also weaken authenticity, reduce emotional connection, and intensify privacy concerns. These tensions become more visible in GenAI contexts because AI is no longer only supporting decisions behind the scenes. It is increasingly involved in producing the very content consumers see and evaluate.

Several gaps remain despite the rapid growth of this field. First, the literature is still fragmented across different applications such as service automation, personalised advertising, disclosure labels, social media engagement, and AI-generated visual or textual content. Second, many studies focus on one variable at a time, such as trust, authenticity, or privacy, without fully explaining how these factors interact in shaping consumer responses. Third, while some empirical studies report positive outcomes for AI when usefulness and convenience are salient, others report negative outcomes when authenticity, emotional sensitivity, or privacy become more important. These differences are often context-specific, but the literature does not yet integrate them clearly enough into one consumer-centred explanation. Finally, broad conceptual discussions of AI and GenAI in marketing often sit separately from narrower empirical studies, making it difficult to construct an integrated account of when AI is perceived positively and when it is resisted [7, 17, 18]. This dissertation addresses these gaps by conducting a thematic literature review focused on the key dimensions through which consumers interpret AI and GenAI in marketing environments.

2.7 Summary

This chapter has reviewed the literature on AI and GenAI in marketing from a consumer perception perspective. It has shown that AI in marketing cannot be assessed only in terms of efficiency or innovation. Consumer responses are shaped by a more complex set of perceptual dimensions, especially trust, authenticity, personalisation, privacy concerns, and behavioural intention. These themes provide the conceptual basis for the methodology and thematic analysis developed in the following chapters.

3 Methodology

3.1 Research Approach

This dissertation adopts a qualitative research approach based on a literature review. Rather than collecting primary data, the study analyses existing academic literature on artificial intelligence (AI), generative artificial intelligence (GenAI), and consumer perception in marketing. This approach is appropriate because the research problem is conceptual and integrative in nature. The aim is not to test one causal relationship statistically, but to synthesise how the literature explains recurring patterns in consumer responses to AI-mediated marketing.

3.2 Research Design

The study uses a thematic literature review design. Thematic analysis was selected because the literature in this area is dispersed across different marketing applications, including personalised advertising, AI-generated content, recommendation systems, chatbots, and social media interactions. A thematic approach allows these studies to be brought together through recurring dimensions of consumer perception rather than through one single application area.

Importantly, this dissertation does not treat positive and negative findings as simple contradictions. Instead, it approaches the literature on the assumption that AI in marketing often generates conditional outcomes and recurring trade-offs. For example, the same AI application may improve relevance and convenience while also increasing privacy concern or reducing perceived authenticity. The methodological aim is therefore to identify not only the main themes in the literature, but also the mechanisms and contextual conditions through which those themes shape consumer responses.

3.3 Data Collection

The data for this study consists of peer-reviewed academic journal articles related to AI in marketing and consumer perception. Relevant literature was identified through academic databases such as Google Scholar, Scopus, and ScienceDirect. These databases were used because they provide broad coverage of marketing, consumer behaviour, and technology-focused research.

Keywords used for searching include:

- Artificial intelligence in marketing
- Generative AI advertising
- Consumer perception AI
- AI-generated content
- AI trust, authenticity, and privacy

Additional searches combined these terms with related concepts such as recommendation systems, disclosure labels, chatbots, virtual influencers, consumer behaviour, and purchase intention. Backward and forward citation tracing was also used where necessary to identify influential studies repeatedly referenced in the field.

Because this dissertation is a thematic literature review rather than a formal systematic review, the search process was iterative rather than fully protocol-driven. Nevertheless, the narrowing process was made explicit. The final corpus used for thematic analysis comprised 26 peer-reviewed journal articles. In addition, two industry reports were retained only for contextual framing in the introduction and background chapter, not as part of the core academic evidence base. Sources were narrowed through staged screening for relevance at title, abstract, and full-text level, with conceptual fit to the research questions taking priority over exhaustive coverage.

3.4 Data Selection Criteria

To ensure relevance and quality, the following criteria were applied:

- Peer-reviewed journal articles
- Primarily published between 2020 and 2026, with earlier studies included where they were conceptually foundational
- A clear focus on AI or generative AI in marketing or advertising contexts

- Direct discussion of consumer perception or closely related concepts such as trust, authenticity, privacy concern, relevance, usefulness, engagement, or purchase intention
- Sufficient conceptual or empirical relevance to the dissertation's research questions

Studies were excluded where they focused mainly on technical system performance, organisational adoption without a consumer dimension, or purely conceptual claims with limited relevance to marketing communication or consumer interpretation. The final body of literature was selected on the basis of relevance and analytical value rather than on a fixed numerical target alone.

In practice, the screening logic moved through three broad stages. First, an initial search and citation-chasing stage produced a broad candidate pool centred on AI, GenAI, consumer response, and digital marketing communication. Second, title and abstract screening removed studies that were clearly technical, organisational, or outside consumer-facing communication contexts. Third, full-text reading was used to retain only those studies that made a direct conceptual or empirical contribution to one or more of the dissertation's core themes. While an exact PRISMA-style record count was not maintained at the start of the project, the final included corpus and the inclusion logic are reported here to make the selection process more transparent.

3.5 Thematic Analysis

Thematic analysis was used to identify key themes across the selected literature and to explain how those themes operate. The process followed several stages:

1. Familiarisation with the literature through repeated reading of selected articles
2. Initial coding of recurring concepts, claims, and findings
3. Grouping of codes into broader perceptual themes
4. Comparison of studies to identify recurring trade-offs, mechanisms, and contextual conditions
5. Refinement of themes into an integrated analytical framework

The coding process was guided by recurrent concepts in the literature, including trust, authenticity, personalisation, privacy concerns, relevance, usefulness, emotional response, and behavioural intention. To structure these concepts more systematically, the dissertation uses Puntoni et al.'s experiential framework as its primary coding lens [18]. Specifically, the literature was coded across four broad experience domains: data capture, classification, delegation, and social interaction. These domains were then translated into four consumer-perception tensions that recur across the reviewed studies:

- Data capture: served versus exploited
- Classification: understood versus misunderstood
- Delegation: empowered versus replaced
- Social interaction: connected versus alienated

These tensions were not treated as isolated variables or formal statistical moderators. Instead, they were used as sensitising themes that helped interpret how consumer responses to AI are produced across different contexts. A fuller coding table showing the relationship between theme groups, key constructs, supporting studies, and interpretive tensions is provided in the Appendix.

Trust requires a further methodological clarification because its role changes slightly between the literature review and the findings chapter. During the initial review stage, trust was retained as a visible thematic heading because it appears repeatedly in prior studies as an explicit construct and often functions as a reported predictor, mediator, or outcome. During the later stage of thematic refinement, however, trust was repositioned as a cross-cutting mechanism because the coding process showed that it emerged across all four experience domains rather than functioning as one self-contained category. This shift was therefore not an ad hoc contradiction in the analysis, but the result of moving from article-level topic coding to cross-study thematic synthesis.

The final analytical framework focuses on three levels of interpretation. First, it identifies the main dimensions through which AI shapes consumer perception. Second, it examines the mechanisms that explain why these effects occur, such as perceived relevance, emotional sincerity, opacity, intrusiveness, or loss of autonomy. Third, it considers the conditions under which responses become more positive or more negative, including AI visibility, message type, product context, and the extent of data use.

To avoid conceptual ambiguity, this dissertation distinguishes between three related but non-identical levels. The first is interpretive appraisal, which refers to how consumers cognitively judge AI as useful, credible, appropriate, intrusive, or fair. The second is affective response, which refers to feelings such as trust, unease, alienation, curiosity, or moral discomfort. The third is behavioural outcome, which refers to consequences such as engagement, ad avoidance, purchase intention, reliance, or resistance. The term consumer perception is therefore used as an umbrella concept that refers most directly to interpretive appraisal and affective response, while behavioural outcomes are treated as consequences shaped by those evaluations rather than as perception itself.

3.6 Limitations

This study has several limitations. First, it relies on secondary data and is therefore dependent on the scope, quality, and contexts of existing published studies. Second, literature selection in a thematic review inevitably involves interpretive judgement, which may shape how themes and mechanisms are prioritised. Third, the field is rapidly developing, particularly in relation to GenAI, meaning that current conclusions should be understood as provisional rather than final. A further limitation is that the evidence base is not evenly distributed across AI types: studies of predictive AI remain more prominent in areas such as personalisation, recommendation, and data-driven targeting, whereas GenAI evidence is more concentrated in visible content creation, disclosure, and authorship-related contexts. Finally, because the dissertation synthesises studies conducted across different markets, platforms, and product categories, some findings may not generalise equally across all consumer contexts.

3.7 Summary

This chapter has outlined the qualitative literature review approach, the selection of academic sources, and the thematic analysis process used in the study. It has also clarified that the purpose of the analysis is not simply to catalogue positive and negative findings, but to explain the mechanisms and conditions through which AI influences consumer perception in marketing. The next chapter presents the findings and discussion developed from this analytical framework.

4 Findings and Discussion

4.1 Introduction

This chapter presents the findings from the thematic analysis of the selected literature and discusses their implications for understanding consumer perception in AI-driven marketing. Rather than treating positive and negative outcomes as contradictory findings, the discussion interprets them as recurring trade-offs and conditional effects. Following Puntoni et al.'s experiential framework [18], the chapter is organised around four recurring tensions: data capture, classification, delegation, and social interaction. This structure makes it possible to show not only what patterns appear in the literature, but also why those patterns arise and under what conditions consumer responses become more favourable or more resistant.

For clarity, the analysis distinguishes between three levels that are closely related but not identical: interpretive appraisal, affective response, and behavioural outcome. In the chapter, consumer perception refers most directly to the first two levels, that is, how consumers make sense of AI and how they feel about it. Behavioural outcomes such as engagement, purchase intention, avoidance, or resistance are treated as consequences shaped by those perceptions rather than as perception itself.

4.2 Interpreting the Four Tensions

The four thematic groups identified in this dissertation should not be understood as four separate lists of positive and negative findings. Rather, each of them captures a recurring tension in how consumers interpret AI in marketing. These tensions help explain why consumer perception of AI is rarely stable or one-dimensional. Across the literature, consumers often value what AI provides while simultaneously resisting the conditions through which those benefits are delivered.

The first tension is between being served and being exploited. AI-driven marketing can make the consumer experience more relevant, efficient, and convenient, but these benefits depend on extensive data capture. As a result, consumers are often drawn into a privacy calculus in which functional value is weighed against surveillance, loss of control, and mistrust. The

second tension is between being understood and being misunderstood. AI classification and personalisation can make consumers feel recognised, yet the same systems may also make them feel stereotyped, reduced to a category, or stripped of individuality.

The third tension is between being empowered and being replaced. Delegating tasks to AI can reduce effort and improve efficiency, but it can also undermine autonomy, self-efficacy, and the sense of authorship over one's own choices. The fourth tension is between being connected and being alienated. AI-mediated communication can create speed, responsiveness, and even novelty, but it may also feel emotionally empty or morally inappropriate when consumers expect genuine empathy, warmth, or human sincerity.

Taken together, these tensions suggest that the literature does not simply contain contradictory findings. Instead, it shows that consumer responses to AI are shaped by recurring trade-offs. AI is more likely to be perceived positively when it delivers practical value without violating expectations of privacy, individuality, autonomy, or authenticity. It is more likely to be perceived negatively when those benefits come at too high a psychological, emotional, or ethical cost. The four sections that follow develop this argument in greater detail.

4.3 Overview of the Coding Framework

The review suggests that consumer perception of AI is shaped by four recurring tensions:

- Data capture: served versus exploited
- Classification: understood versus misunderstood
- Delegation: empowered versus replaced
- Social interaction: connected versus alienated

These tensions do not operate independently. Across the literature, consumer responses to AI are shaped by the interaction between perceived benefits such as convenience, relevance, efficiency, and responsiveness, and perceived costs such as intrusiveness, opacity, stereotyping, loss of autonomy, and perceived inauthenticity. The findings therefore suggest that AI in marketing should be understood through a relational and interpretive lens, not simply through a performance lens. Trust, authenticity, privacy concern, and behavioural intention are best understood as outcomes or mechanisms distributed across these four broader experience domains rather than as fully separate themes.

4.4 Data Capture and Privacy: Served versus Exploited

The first thematic tension concerns data capture. AI-driven marketing depends heavily on consumer data in order to personalise recommendations, optimise advertising, and support real-time interaction. From the consumer perspective, this can create a positive experience when data use produces relevance, convenience, and reduced search effort. In such cases, consumers may feel served by AI because the system appears to anticipate their needs and improve the efficiency of the market experience [14, 19].

However, the same process can produce a very different perception when data collection feels opaque, excessive, or difficult to control. Aguirre et al. show that personalised advertising performs better when information collection is overt, but that covert collection heightens perceived vulnerability and reduces effectiveness [22]. Hardcastle et al. similarly show that AI-driven personalised journeys can create frustration and resentment when consumers feel persistently monitored or when stored data seems overly extensive [23]. Boerman et al. likewise show that targeted digital advertising can quickly shift from relevance to perceived intrusiveness when consumers become more aware of how data is being collected and used [21]. Martin and Murphy provide broader conceptual support by showing that privacy in marketing is not a peripheral issue but a structural concern tied to trust, legitimacy, and power in data use [24]. More broadly, Davenport et al. and Kumar et al. identify privacy and data governance as structural tensions in AI-enabled marketing [2, 16]. Ozturkcan and Bozdağ extend this concern by showing how opaque practices and exaggerated claims about AI can generate a broader cycle of mistrust, in which AI washing is followed by consumer backlash or “AI booing” [25]. Srivastava and Gurme further show that privacy concern and perceived surveillance can directly weaken trust even when personalisation remains technically effective [15]. These studies suggest that the issue is not simply whether firms use data, but whether that use is interpreted as proportionate, transparent, and legitimate.

What consumers perceive in this theme is a tension between being efficiently served and being quietly exploited. Consumers value tailored relevance, convenience, and reduced search effort, but they also become highly sensitive to signs that their personal data is being continuously monitored, stored, or abstracted into computational profiles without meaningful consent. This is why the same AI system can be interpreted either as useful service or as invasive surveillance.

Why this perception occurs can be explained through privacy calculus and perceived vulnerability. Consumers weigh the practical benefits of personalisation against the costs of dataveillance, loss of privacy, and reduced control. This tension becomes especially intense

when data collection is covert or when the logic of the system remains opaque. It is further exacerbated when firms overstate their AI sophistication or ethical responsibility, because AI washing inflates expectations that are then violated when the system appears exploitative rather than trustworthy. The underlying mechanism is therefore not merely data use itself, but the interpretation of that data use as either legitimate service or covert extraction.

The main behavioural outcomes affected by this tension are trust, click-through intention, ad avoidance, purchase intention, and longer-term willingness to remain engaged with the customer journey. Transparent and proportionate data practices are more likely to sustain trust and purchase intention, whereas opaque or excessive data capture is more likely to produce avoidance, reactance, and, in more severe cases, public backlash or disengagement [14, 15]. The evidence for this theme is especially strong because similar patterns appear across experiments, surveys, and review-based studies. The tension is therefore best understood not as a contradiction in the literature, but as a privacy-personalisation trade-off built into the data capture experience itself.

4.5 Classification and Personalisation: Understood versus Misunderstood

The second thematic tension concerns classification. AI systems sort consumers into categories, infer their likely preferences, and use those classifications to generate personalisation. When this process works well, consumers may feel understood because the message, recommendation, or service appears relevant to their actual needs. This is why personalisation is often associated with favourable consumer response in the literature [1, 2, 14].

Yet the classification process can also make consumers feel misunderstood. Puntoni et al. suggest that AI can produce experiences of misunderstanding when it reduces people to narrow inferred categories or makes judgements that fail to reflect their full identity and intention [18]. Hou et al. provide particularly direct evidence for this problem by showing that consumers often infer that AI-generated personalised solutions are shared with a larger group of people and are therefore less unique [26]. Starke et al. further show that perceptions of algorithmic fairness are highly sensitive to context, procedure, and outcome, which helps explain why classification can also trigger concerns about stereotyping, bias, and unequal treatment [27]. Longoni et al. add an important boundary condition by showing that consumers become especially resistant when AI is perceived to lack the judgement needed for high-stakes or personally nuanced decisions [28]. Even where personalisation is technically accurate, it may still feel mechanical or impersonal if the consumer believes that many others are being treated in essentially the same way. The issue is therefore not only

predictive success, but whether the classification feels sufficiently individual and context-sensitive.

What consumers perceive in this theme is a tension between being recognised and being reduced. On the positive side, AI classification can make consumers feel understood when the system appears to anticipate needs and provide relevant solutions. On the negative side, consumers may feel misunderstood when they suspect that the system relies on stereotypes, rigid templates, or overgeneralised profiles rather than on an appreciation of their individual complexity.

Why this perception occurs can be explained through uniqueness neglect and fairness concerns. Because consumers often attribute low experiential sensitivity to AI, they assume that algorithmic recommendations are generated through broad rules rather than genuine insight. This leads them to infer a larger expected group size, meaning that the personalised solution does not feel rare, exclusive, or truly individual. Where biased or historically skewed data are involved, this perception becomes more severe because classification can also be interpreted as unfair or discriminatory. The central mechanism is therefore not simply whether the recommendation is personalised, but whether the underlying classification appears psychologically sensitive and procedurally fair.

The behavioural outcomes affected by this theme include adoption intention, willingness to rely on AI advice, willingness to pay for personalised solutions, and acceptance of AI providers in higher-stakes contexts. When classification produces a convincing sense of understanding, AI is more likely to be accepted as useful and competent. When it produces uniqueness neglect or fairness concerns, consumers become more reluctant to trust the recommendation, less willing to adopt it, and more likely to reject AI providers in contexts where individual nuance matters. The evidence base here is moderate to strong: the overall pattern is clear, even though fairness perceptions remain highly context-sensitive. The classification experience therefore introduces a subtle but important trade-off between precision and reductionism.

4.6 Task Delegation and Autonomy: Empowered versus Replaced

The third thematic tension concerns delegation. Consumers often rely on AI to assist with search, recommendation, customer service, and decision support. In many contexts this delegation is experienced positively because it reduces effort, saves time, and makes the market environment easier to navigate. Ameen et al. show that AI-enabled services can improve customer experience when they are perceived as competent, convenient, and supportive of the consumer's goals [19]. In these situations, delegation is experienced as

empowerment rather than dependence.

At the same time, delegation can generate a loss of autonomy. Puntoni et al. argue that AI may create tension when consumers feel that judgement, choice, or effort has been transferred too far away from the self [18]. Hardcastle et al. suggest that this tension intensifies when predictive journeys repeatedly direct consumers toward similar content or options, effectively narrowing what they see and reinforcing a sense of constrained choice [23]. Davenport et al. similarly note that AI-driven micro-targeting can reduce search costs while also undermining the consumer's perceived freedom of choice [16]. This becomes especially relevant when AI does not merely support a decision but appears to narrow the range of options, shape preferences too strongly, or replace forms of action that consumers associate with agency, creativity, or competence. In such cases, the consumer no longer experiences AI as assistance but as partial replacement.

What consumers perceive in this theme is a tension between empowerment and displacement. AI can feel empowering when it removes mundane effort and helps consumers navigate large amounts of information more efficiently. Yet the same delegation can feel threatening when consumers sense that control is shifting away from them, especially in tasks that involve subjectivity, judgement, or identity.

Why this perception occurs can be explained through perceived sacrifice and loss of autonomy. AI systems can guide choices so heavily that consumers no longer experience themselves as the true authors of their decisions. This becomes more serious when algorithmic curation narrows visible options or creates filter bubbles, because the convenience of delegation starts to look like a loss of freedom. In identity-relevant activities, delegation may also feel like the loss of an accomplishment that consumers believe should belong to them rather than to a machine. The core mechanism here is the conversion of convenience into psychological threat when support turns into substitution.

The behavioural outcomes associated with this theme include resistance to recommendations, frustration, reduced reliance on AI support, and psychological reactance. In some cases, consumers may deliberately ignore or reject algorithmic guidance simply to reassert their independence. Although the empirical base on long-term autonomy loss is still developing, the existing evidence consistently suggests that delegation is most positively perceived when it remains supportive and reversible, and most negatively perceived when it appears to replace rather than augment the consumer's own agency.

4.7 Social Interaction and Authenticity: Connected versus Alienated

The fourth thematic tension concerns social interaction. Consumers increasingly encounter AI in socially framed contexts, such as chatbots, AI-generated advertising, virtual assistants, and other forms of automated brand communication. These interactions can create a sense of connection when they are responsive, clear, and useful. Consumers may appreciate instant assistance, reduced waiting time, and the novelty of AI-mediated brand interaction, especially in factual or utilitarian settings [17,19].

However, this same social experience can become alienating when AI communication feels emotionally empty, morally inappropriate, or insufficiently human. This is where authenticity becomes especially important. As discussed in the literature review, authenticity in this dissertation has two dimensions: source authenticity and emotional authenticity. Source authenticity concerns whether the message is perceived as genuinely human-authored or AI-authored [9]. Emotional authenticity concerns whether the communication feels sincere, empathetic, and genuinely intended [20]. These two dimensions help explain why consumers may resist AI not simply because it is artificial, but because it disrupts expected forms of warmth, empathy, or symbolic meaning.

What consumers perceive in this theme is a dual response: AI can feel novel, responsive, and efficient, yet also inauthentic, emotionally empty, or even deceptive. Consumers may enjoy the speed and availability of AI-mediated communication, but they remain aware that machines do not possess genuine internal emotional states. Kirk and Givi show that this can produce an AI-authorship effect in which the perceived lack of real feeling weakens response to emotionally expressive communication [9]. Shi and Jiang similarly show that AI disclosure can increase novelty while simultaneously reducing authenticity [10]. This creates a tension between the convenience of connection and the discomfort of artificiality.

Why this perception occurs can be explained through two competing mechanisms. The first is perceived novelty, which can stimulate curiosity and engagement because AI signals innovation. The second, often more powerful, is perceived inauthenticity. When AI attempts to simulate empathy, warmth, or emotional sincerity, consumers may judge the interaction as fake because the expression is not backed by genuine feeling. Arango et al. show how this problem can reduce empathy and donation intention in charitable advertising [20], while Belanche et al. show that resistance becomes stronger in hedonic or high-involvement settings where realism and genuine human presence matter more [12]. In some contexts this mismatch generates moral disgust, especially when AI is used in emotionally charged or morally sensitive communication. The effect is therefore strongly context-dependent: utilitarian tasks invite more acceptance, whereas hedonic or high-involvement experiences

intensify the need for human touch and authentic presence.

The behavioural outcomes associated with this theme include short-term engagement through novelty, but also reduced word-of-mouth, lower loyalty, weaker donation intention, and stronger resistance when authenticity concerns dominate. In medical, charitable, and emotionally expressive settings, alienation can become especially strong because consumers fear that their individuality, seriousness, or emotional needs are not being genuinely recognised [28]. The evidence for this theme is strong across multiple experiments and applied contexts. Social interaction with AI is therefore best understood as a double-edged mechanism that can create both connection and alienation depending on how the communication is framed and where it occurs.

4.8 How Trust, Authenticity, and Behavioural Intention Fit Across the Four Themes

The analysis suggests that trust, authenticity, and behavioural intention should not be treated as isolated findings detached from the broader experience structure. In the literature review, trust was discussed as a distinct theme because it is one of the most visible concepts in prior studies and often appears as an explicit dependent or mediating variable. In the findings, however, trust is repositioned as a cross-cutting mechanism because the thematic synthesis shows that it is produced through several different experience domains rather than operating as one separate empirical topic. Trust is shaped across data capture, classification, delegation, and social interaction, which is why it is analytically more useful here as an evaluative mechanism than as a fifth standalone theme.

Authenticity is primarily concentrated in the social interaction theme, but it also interacts with trust because doubts about authorship and emotional sincerity can weaken the credibility of AI-mediated communication. Behavioural intention, by contrast, is best understood as an outcome of how the four tensions are resolved. Consumers are more likely to engage, purchase, or respond favourably when they feel served, understood, empowered, and connected. They are less likely to do so when they feel exploited, misunderstood, replaced, or alienated.

4.9 How the Four Tensions Interact

The four tensions identified in this dissertation should not be treated as independent blocks. Although they are presented separately for analytical clarity, the literature suggests that they often reinforce or condition one another. In practice, consumers do not encounter data capture, classification, delegation, and social interaction as isolated dimensions. They

encounter them together within the same digital experience, which means that one tension can intensify or moderate another.

The clearest interaction appears between data capture and classification. Consumers are less likely to feel understood by personalisation when they are already uneasy about how their data was collected. In this sense, privacy concern can weaken the positive effects of relevance, because the same recommendation may be interpreted less as helpful recognition and more as intrusive profiling. A similar interaction exists between classification and delegation. When AI appears to rely on narrow categories or generic rules, consumers are less willing to delegate judgement to it, especially in contexts that require nuance, taste, or personal identity. Misunderstanding therefore makes replacement feel more threatening.

Data capture also interacts with social interaction. AI-mediated communication may already feel emotionally thin or inauthentic, but that concern can become stronger when the message is also understood as the product of hidden surveillance or excessive profiling. In other words, perceived intrusion can make perceived inauthenticity more severe. Conversely, when AI use is transparent, proportionate, and clearly bounded, consumers may be more willing to accept novelty or limited automation in communication.

Taken together, these patterns suggest that the four tensions form an interpretive system rather than a simple list. Data concerns can shape whether personalisation feels caring or exploitative. Classification can shape whether delegation feels supportive or substitutive. Social interaction can amplify how questions of trust, privacy, and authenticity are ultimately resolved. This interactional reading strengthens the framework by showing that consumer perception is shaped not only by the presence of individual tensions, but by how those tensions accumulate and condition each other within a single AI-mediated encounter.

4.10 Conditions That Shape More Positive or Negative Responses

The literature suggests that consumer responses to AI become more positive or more negative under identifiable conditions. One important condition is the visibility of AI. When AI works in the background to improve convenience or targeting, it is often evaluated more positively. When it becomes highly visible as the creator of emotionally meaningful or identity-relevant communication, consumers become more alert to questions of authenticity and trust.

Another condition is message type. AI appears to be more favourably received when the message is informational, functional, or service-oriented than when it is expected to convey emotional sincerity, artistic originality, or social meaning. Product and context also matter.

High-involvement, hedonic, and symbolic categories tend to intensify concerns about authenticity and human judgement, whereas lower-stakes or convenience-driven contexts often make efficiency benefits more salient [12, 20].

Data practice is a further condition. Consumers are more willing to accept AI when personalisation feels proportionate and understandable. They become more resistant when the same personalisation feels opaque, excessive, or difficult to control [15, 21]. These patterns suggest that the consumer response to AI is not random. It is shaped by recurring contextual conditions that determine whether the benefits of AI outweigh its relational and ethical costs.

4.11 Implications of the Findings

For researchers, the findings suggest that future work should move beyond simply identifying whether AI has positive or negative effects. What matters more is specifying the mechanisms and conditions through which those effects occur. This would allow future studies to build a more cumulative understanding of consumer response to AI and GenAI across different marketing settings.

For practitioners, the main implication is that AI should not be managed only as an efficiency tool. It must also be managed as a communicative and relational presence. The success of AI in marketing depends on whether it is used in ways that preserve trust, respect privacy expectations, and match the symbolic demands of the message context. In practice, this means that high precision, fast content generation, or automated interaction are not enough on their own. They must be aligned with what consumers interpret as appropriate, credible, and worth engaging with.

4.12 Summary

This chapter has shown that consumer perception of AI in marketing is shaped by recurring trade-offs rather than by simple contradiction. Organised through the four tensions of data capture, classification, delegation, and social interaction, the analysis suggests that AI tends to generate more positive responses when consumers feel served, understood, empowered, and connected, but more negative responses when they feel exploited, misunderstood, replaced, or alienated. These outcomes are produced through identifiable mechanisms and shaped by identifiable conditions, including AI visibility, message type, emotional expectation, and data practice. The next chapter concludes the dissertation by summarising these findings and outlining their broader implications.

5 Conclusion

5.1 Introduction

This chapter summarises the key findings of the study, discusses theoretical and practical implications, and outlines limitations and directions for future research.

5.2 Summary of Findings

This dissertation examined how artificial intelligence (AI) and generative artificial intelligence (GenAI) influence consumer perception in marketing. The study argued that consumer perception of AI is not adequately explained by describing AI as either beneficial or harmful in general terms. Instead, the literature suggests that consumer responses are shaped by recurring tensions that emerge when AI offers practical value while simultaneously creating psychological, emotional, and ethical concerns.

Using a thematic literature review informed by Puntoni et al.'s experiential framework, the dissertation identified four overarching tensions: data capture, classification, delegation, and social interaction. These were interpreted through the consumer-perception tensions of served versus exploited, understood versus misunderstood, empowered versus replaced, and connected versus alienated. Taken together, these tensions show that consumer responses to AI in marketing are conditional rather than uniform.

The first key finding is that AI tends to generate more favourable responses when data use is perceived as proportionate, transparent, and valuable. Consumers are more likely to accept AI when it improves relevance, convenience, and efficiency without making them feel monitored or exploited. The second key finding is that personalisation does not automatically create a sense of individuality. AI can make consumers feel understood, but it can also make them feel reduced to a category when recommendations seem generic, stereotypical, or insufficiently responsive to personal context.

The third key finding is that delegation produces both relief and resistance. AI is often welcomed when it reduces effort and supports decision-making, but it becomes more

problematic when consumers feel that control, agency, or self-efficacy is being displaced. The fourth key finding is that authenticity becomes especially important when AI moves from back-end optimisation into visible content generation and social interaction. Consumers may accept AI in factual or utilitarian contexts, yet become more resistant in emotional, hedonic, or morally sensitive contexts where they expect human sincerity, empathy, or creative intention.

Across all four tensions, trust emerged as a cross-cutting mechanism rather than a standalone theme. Trust increases when AI is perceived as useful, legitimate, and contextually appropriate. It declines when AI appears opaque, manipulative, overly visible, or emotionally inauthentic. Behavioural outcomes such as engagement and purchase intention are therefore best understood as consequences of how these tensions are resolved in specific contexts.

5.3 Theoretical Implications

This study contributes to the literature in several ways. First, it offers a more precise conceptualisation of consumer perception in AI-mediated marketing by treating perception as an interpretive process that combines cognitive evaluation, affective response, and behavioural consequence. This helps move the discussion beyond narrower adoption-oriented approaches that are less suited to explaining visible and socially meaningful AI encounters in marketing.

Second, the dissertation contributes by showing that many apparent contradictions in the literature are more accurately understood as trade-offs or conditional effects. This matters theoretically because it shifts the focus from asking whether AI is positively or negatively perceived to asking under what conditions certain benefits are accepted and under what conditions those same benefits are reinterpreted as risks.

Third, the study demonstrates the value of Puntoni et al.'s experiential framework for organising the AI-consumer literature. By translating data capture, classification, delegation, and social interaction into four consumer-perception tensions, the dissertation provides a more integrated framework for connecting privacy, trust, authenticity, autonomy, and behavioural intention. This offers a clearer structure for future literature reviews and empirical research.

Finally, the dissertation extends existing discussion of GenAI by distinguishing between source authenticity and emotional authenticity. This distinction is theoretically useful because it clarifies that consumer resistance to AI-generated communication is not only about whether a message is AI-authored, but also about whether it feels emotionally sincere and contextually appropriate.

5.4 Practical Implications

From a practical perspective, the findings suggest that marketers should carefully balance the use of AI technologies with consumer expectations.

First, firms should not assume that technological sophistication automatically improves consumer response. AI is most effective when the value it creates is clearly legible to the consumer. This means that personalisation, recommendation, and service automation should be designed to feel supportive rather than invasive.

Second, transparency and proportionality in data practice are essential. The findings suggest that consumers are more willing to accept AI when data use feels legitimate and understandable. Marketers should therefore reduce opaque targeting practices, avoid overstating AI capability, and ensure that data capture remains proportionate to the benefit being delivered.

Third, firms should be cautious when deploying AI in emotionally expressive, hedonic, or morally sensitive communication. In such contexts, efficiency alone is unlikely to secure positive consumer response. Brands need to consider whether AI-generated or AI-disclosed communication may weaken authenticity, trust, or emotional connection.

Fourth, AI should be positioned as augmenting rather than replacing consumer agency. Systems that help consumers navigate choice, gather information, or receive service may be positively received, but systems that appear to narrow options, automate judgement too aggressively, or remove too much human presence may generate resistance.

5.5 Limitations

This study has several limitations. First, it is based entirely on secondary research and therefore depends on the scope, quality, and contextual diversity of existing studies. The dissertation does not provide new primary empirical evidence and cannot directly test the mechanisms proposed in the analysis.

Second, the literature review is interpretive in nature. Although the thematic framework provides coherence, the selection and synthesis of themes inevitably involve judgement. Different reviewers might emphasise somewhat different mechanisms, contextual conditions, or boundaries between themes.

Third, the field is developing rapidly, especially in relation to GenAI. As a result, some of the conclusions drawn here should be understood as provisional and open to revision as new empirical evidence emerges. A further limitation is that the dissertation title includes both AI and GenAI, while the available evidence is not evenly distributed across those two

domains. Research on predictive AI remains more dominant in areas such as recommendation systems, personalised advertising, and data-driven targeting, whereas GenAI-specific evidence is much more concentrated in visible content generation, disclosure, and authorship-related contexts. This means that the four tensions identified in the study apply across AI more broadly, but the GenAI contribution is currently most developed in the social interaction and authenticity literature. Finally, the reviewed literature spans different platforms, markets, product categories, and cultural settings, which means that the general patterns identified in this study may not apply equally across all contexts.

5.6 Future Research

Future research should build on this dissertation in at least four directions. First, empirical studies are needed to test the mechanisms identified here, especially the roles of trust, perceived exploitation, perceived understanding, perceived control, and perceived human presence in shaping behavioural outcomes.

Second, future research should examine the conditions under which AI is more likely to be accepted or resisted. In particular, product type, involvement level, emotional intensity, cultural context, and disclosure strategy appear especially important and merit more systematic investigation.

Third, more research is needed on the distinction between source authenticity and emotional authenticity in GenAI marketing. This would help clarify when consumers object primarily to AI authorship itself and when they object more strongly to the emotional inadequacy of the resulting communication.

Fourth, future work should examine differences across consumer groups, markets, and levels of AI literacy. The literature suggests that not all consumers interpret AI in the same way, and a more differentiated understanding of these variations would strengthen both theory and practice.

Fifth, future research should extend GenAI inquiry beyond authorship and visible content effects. At present, GenAI is most often examined in relation to disclosure, authenticity, and AI-generated communication, while less is known about how GenAI may reshape the other three tensions through data capture, classification, and delegation. More work is therefore needed to determine whether GenAI introduces genuinely new perceptual mechanisms in those areas or intensifies dynamics already documented in predictive AI research.

5.7 Final Remarks

Artificial intelligence and generative AI are transforming the marketing landscape. While they offer significant opportunities for improving efficiency and personalisation, they also introduce new challenges for consumer perception.

This dissertation has argued that the central question is not whether AI in marketing is simply good or bad for consumers. The more important question is how consumers interpret the trade-offs AI creates, and under what conditions those trade-offs become acceptable or unacceptable. By showing that consumer perception is shaped by the tensions of being served or exploited, understood or misunderstood, empowered or replaced, and connected or alienated, the study provides a clearer framework for understanding the evolving relationship between AI, marketing, and consumer experience.

Bibliography

- [1] A. Haleem, M. Javaid, M. A. Qadri, R. P. Singh, and R. Suman, "Artificial intelligence (AI) applications for marketing: A literature-based study," *International Journal of Intelligent Networks*, vol. 3, pp. 119–132, 2022.
- [2] V. Kumar, A. R. Ashraf, and W. Nadeem, "AI-powered marketing: What, where, and how?" *International Journal of Information Management*, vol. 77, p. 102783, 2024.
- [3] R. Jain and A. Kumar, "Artificial intelligence in marketing: Two decades review," *NMIMS Management Review*, vol. 32, no. 2, pp. 75–83, 2024.
- [4] A. Singla, A. Sukharevsky, B. Hall, L. Yee, M. Chui, and T. Balakrishnan, "The state of AI in 2025: Agents, innovation, and transformation," McKinsey & Company, Industry report, Nov. 2025.
- [5] Deloitte AI Institute, "From concept to consumer: Unlocking GenAI value in retail and consumer products," Deloitte, Industry report, 2025.
- [6] E. Hermann and S. Puntoni, "Artificial intelligence and consumer behavior: From predictive to generative AI," *Journal of Business Research*, vol. 180, p. 114720, 2024.
- [7] V. O. Ubal, M. Lisjak, and M. Mende, "Cracking the consumers' code: A framework for understanding the artificial intelligence–consumer interface," *Current Opinion in Psychology*, vol. 58, p. 101832, 2024.
- [8] M. R. Mohd Amin, A. Asbi, V. M. Sivakumaran, J. Kim, and E. Septiarini, "Artificial intelligence (AI) adoption in marketing strategies: Navigating the present and shaping the future business landscape," *Social Sciences & Humanities Open*, vol. 12, p. 102048, 2025.
- [9] C. P. Kirk and J. Givi, "The AI-authorship effect: Understanding authenticity, moral disgust, and consumer responses to AI-generated marketing communications," *Journal of Business Research*, vol. 186, p. 114984, 2025.
- [10] Y. Shi and Z. Jiang, "Consumer responses to AI disclosure labels: The role of novelty and authenticity," *Sage Open*, vol. 16, no. 1, p. 21582440261417793, 2026.

- [11] K. Israfilzade, "AI-generated versus human-created advertising: Effects on consumer trust and purchase intent," *Equilibrium. Quarterly Journal of Economics and Economic Policy*, vol. 20, pp. 1301–1337, 2025.
- [12] D. Belanche, S. Ibáñez-Sánchez, P. Jordán, and S. Matas, "Customer reactions to generative AI vs. real images in high-involvement and hedonic services," *International Journal of Information Management*, vol. 85, p. 102954, 2025.
- [13] L. Zhang and C. Hur, "The impact of generative AI images on consumer attitudes in advertising," *Administrative Sciences*, vol. 15, no. 10, p. 395, 2025.
- [14] G. K. An and T. T. A. Ngo, "AI-powered personalized advertising and purchase intention in vietnam's digital landscape: The role of trust, relevance, and usefulness," *Journal of Open Innovation: Technology, Market, and Complexity*, vol. 11, no. 3, p. 100580, 2025.
- [15] R. K. Srivastava and V. Gurme, "Artificial intelligence and consumer behaviour on social media: A study of personalization, trust, and privacy," *Measurement: Digitalization*, vol. 5, p. 100021, 2026.
- [16] T. Davenport, A. Guha, D. Grewal, and T. Bressgott, "How artificial intelligence will change the future of marketing," *Journal of the Academy of Marketing Science*, vol. 48, no. 1, pp. 24–42, 2020.
- [17] H.-L. Chan and T.-M. Choi, "Using generative artificial intelligence (GenAI) in marketing: Development and practices," *Journal of Business Research*, vol. 191, p. 115276, 2025.
- [18] S. Puntoni, R. W. Reczek, M. Giesler, and S. Botti, "Consumers and artificial intelligence: An experiential perspective," *Journal of Marketing*, vol. 85, no. 1, pp. 131–151, 2021.
- [19] N. Ameen, A. Tarhini, A. Reppel, and A. Anand, "Customer experiences in the age of artificial intelligence," *Computers in Human Behavior*, vol. 114, p. 106548, 2021.
- [20] L. Arango, S. P. Singaraju, and O. Niininen, "Consumer responses to AI-generated charitable giving ads," *Journal of Advertising*, vol. 52, no. 4, pp. 486–503, 2023.
- [21] S. C. Boerman, S. Kruikemeier, and F. J. Zuiderveen Borgesius, "Online behavioral advertising: A literature review and research agenda," *Journal of Advertising*, vol. 46, no. 3, pp. 363–376, 2017.
- [22] E. Aguirre, D. Mahr, D. Grewal, K. de Ruyter, and M. Wetzels, "Unraveling the personalization paradox: The effect of information collection and trust-building

- strategies on online advertisement effectiveness," *Journal of Retailing*, vol. 91, no. 1, pp. 34–49, 2015.
- [23] K. Hardcastle, L. Vorster, and D. M. Brown, "Understanding customer responses to AI-Driven personalized journeys: Impacts on the customer experience," *Journal of Advertising*, vol. 54, no. 2, 2025.
- [24] K. D. Martin and P. E. Murphy, "The role of data privacy in marketing," *Journal of the Academy of Marketing Science*, vol. 45, no. 2, pp. 135–155, 2017.
- [25] S. Ozturkcan and A. A. Bozdağ, "Responsible AI in marketing: AI booing and AI washing cycle of AI mistrust," *International Journal of Market Research*, vol. 67, no. 6, pp. 696–722, 2025.
- [26] C. Hou, T. Li, and Y. Gu, "Artificial intelligence versus human providers for personalized solutions? the influence of expected group size and perceived uniqueness on adoption intentions," *Journal of Retailing and Consumer Services*, vol. 88, p. 104544, 2026.
- [27] C. Starke, J. Baleis, B. Keller, and F. Marcinkowski, "Fairness perceptions of algorithmic decision-making: A systematic review of the empirical literature," *Big Data & Society*, vol. 9, no. 2, 2022.
- [28] C. Longoni, A. Bonezzi, and C. K. Morewedge, "Resistance to medical artificial intelligence," *Journal of Consumer Research*, vol. 46, no. 4, pp. 629–650, 2019.

A1 Appendix

A1.1 Coding Framework Overview

This appendix provides a summary of the thematic coding framework used in the dissertation. The coding process was informed by Puntoni et al.'s experiential framework and translated into four overarching consumer-perception tensions. The appendix is included to show how the main analytical categories used in Chapter 3 and Chapter 4 were derived from the literature.

Table A1.1: Summary of the thematic coding framework

Theme group	Puntoni domain	Core tension	Illustrative constructs
Data capture and privacy	Data capture experience	Served versus exploited	Privacy calculus, dataveillance, perceived vulnerability, transparency, trust
Classification and personalisation	Classification experience	Understood versus misunderstood	Personalisation, uniqueness neglect, expected group size, stereotyping, algorithmic fairness
Task delegation and autonomy	Delegation experience	Empowered versus replaced	Self-efficacy, perceived control, perceived sacrifice, autonomy, algorithm aversion
Social interaction and authenticity	Social experience	Connected versus alienated	Source authenticity, emotional authenticity, empathy gap, moral disgust, novelty

A1.2 Coding Table

Table A1 provides a fuller version of the coding logic used in the dissertation. It links the main theme groups to associated concepts, supporting studies, and their interpretive relevance to consumer perception.

Table A1.2: Full thematic coding table

Theme group	Key constructs	Illustrative supporting studies	Interpretive relevance
Data capture and privacy	Privacy calculus, surveillance, dataveillance, trust, transparency, perceived exploitation	Aguirre et al. (2015); Boerman et al. (2017); Martin and Murphy (2017); Hardcastle et al. (2025); An and Ngo (2025); Srivastava and Gurme (2026); Ozturkcan and Bozdağ (2025)	Shows how consumers may value tailored service while simultaneously resisting opaque or excessive data capture
Classification and personalisation	Personalisation, algorithmic classification, uniqueness neglect, expected group size, bias, fairness	Puntoni et al. (2021); Hou et al. (2026); Starke et al. (2022)	Shows that AI may create relevance without necessarily creating a feeling of individuality or fair recognition
Task delegation and autonomy	Delegation, self-efficacy, autonomy, perceived control, replacement, psychological threat	Puntoni et al. (2021); Davenport et al. (2020); Ameen et al. (2021); Hardcastle et al. (2025)	Shows how AI can reduce effort while also narrowing choice and weakening the consumer's sense of agency
Social interaction and authenticity	Source authenticity, emotional authenticity, empathy gap, novelty, alienation, moral disgust	Kirk and Givi (2025); Shi and Jiang (2026); Arango et al. (2023); Belanche et al. (2025); Longoni et al. (2019)	Shows how AI-mediated communication may be accepted in factual contexts but resisted in emotional or high-involvement settings

A1.3 Appendix Note

The appendix coding tables are included to make the analytical process transparent. They are not intended to replace the discussion in the main chapters, but to document how the dissertation moved from literature selection to thematic interpretation.